

01_assignment_perceptron_robertson

February 15, 2025

Assignment 1: Perceptron

CPSC 381/581: Machine Learning

Yale University

Instructor: Alex Wong

Student: Hailey Robertson, hdr22

Prerequisites:

1. Enable Google Colaboratory as an app on your Google Drive account
2. Create a new Google Colab notebook, this will also create a “Colab Notebooks” directory under “MyDrive” i.e.

`/content/drive/MyDrive/Colab Notebooks`

3. Create the following directory structure in your Google Drive

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments`

4. Move the `01_assignment_perceptron_multiclass.ipynb` into

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments`

so that its absolute path is

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments/01_assignment`

In this assignment, you will implement both binary and multiclass perceptron classifiers from scratch. You will test your implementations on the digits dataset from scikit-learn. The assignment is divided into three main parts:

1. Implementing a binary perceptron for digit classification (0 vs 1)
2. Implementing a multiclass perceptron for full digits classification (0-9)
3. Comparing your implementations with scikit-learn’s Perceptron

Submission:

1. Implement all TODOs in the code blocks below.
2. Report your validation and testing scores. For full credit, your testing scores should be higher than 0.9.

BinaryPerceptron:

Best model test accuracy: 0.9861

Scikit-learn Perceptron test accuracy: 1.0000

MultiClassPerceptron:

Best model test accuracy: 0.9611

Scikit-learn Perceptron test accuracy: 0.9361

3. List any collaborators.

Collaborators: N/A

```
[26]: import numpy as np
import sklearn.datasets as skdata
from sklearn.linear_model import Perceptron
import sklearn.metrics as skmetrics
from sklearn.model_selection import train_test_split
import warnings, time
import matplotlib.pyplot as plt

warnings.filterwarnings(action='ignore')
np.random.seed(42)
```

```
[27]: class BinaryPerceptron:
    """
    Implementation of Binary Perceptron
    """

    def __init__(self):
        self.__weights = None

    def __update(self, x, y):
        """
        Update weights for misclassified examples

        Arg(s):
            x : numpy.ndarray
                Feature vector of shape d x 1
            y : int
                Label/target (-1 or 1)
        """

        # DONE: Implement weight update rule for binary perceptron
        # NOTE: x = (d+1, 1) column vector with bias, so select first column
        ↪ for (d+1,) array
            self.__weights[:, 0] += y * x

    def fit(self, x, y, max_iter=100):
        """
        Fit the binary perceptron to training data
        """
```

```

    Arg(s):
        x : numpy.ndarray
            Features of shape  $d \times N$ 
        y : numpy.ndarray
            Labels/targets of shape  $(1 \times N)$  or  $(N \times 1)$ 
        max_iter : int
            Maximum number of iterations
    '''

    n_features, n_samples = x.shape

    # DONE: Initialize weights (including a bias term,  $w_0$ ) as zeros vector
    ↪ with shape  $d+1 \times 1$ 
    self.__weights = np.zeros((n_features + 1, 1))

    # DONE: Append artificial coordinate ( $x_0$ ) to the data
    # NOTE: This makes the shape  $d+1 \times 1$  since bias isn't separate
    x = np.concatenate((np.ones((1, n_samples))), x), axis=0)

    # DONE: Implement training loop
    for _ in range(max_iter):
        n_updates = 0

        # Process each sample
        for n in range(n_samples):

            # DONE: Calculate prediction
            binary_prediction = np.sign(self.__weights.T @ x[:, n])
            # print(binary_prediction, y[n]) # shape (1,1) but single value

            # DONE: Update weights if misclassified
            if binary_prediction != y[n]:
                self.__update(x[:, n], y[n])
                # print(self.__weights)
                n_updates += 1

            # DONE: Break if no updates were made, e.g., check for convergence
            if n_updates == 0:
                break

    def predict(self, x):
        '''
        Make predictions

        Arg(s):
            x : numpy.ndarray

```

```

        Features of shape d x N

    Returns:
        numpy.ndarray : Predicted labels (-1 or 1) of 1 x N
    """

    n_features, n_samples = x.shape

    # DONE: Append artificial coordinate (x0) to the data
    x = np.concatenate((np.ones((1, n_samples))), x), axis=0)

    # DONE: Implement prediction logic
    # NOTE: Was having a lot of trouble getting shapes to match (see notes_
    ↪ above) so ended up doing .flatten() to deal with scalar
    binary_prediction = np.sign(self.__weights.T @ x).flatten()
    return binary_prediction

def score(self, x, y):
    """
    Calculate prediction accuracy

    Arg(s):
        x : numpy.ndarray
            Features of shape d x N
        y : numpy.ndarray
            Labels/targets of shape (1 x N) or (N x 1)

    Returns:
        float: Accuracy score
    """

    # DONE: Implement accuracy calculation
    binary_predictions = self.predict(x)
    binary_accuracy = np.mean(binary_predictions == y.flatten())
    return binary_accuracy

```

```

[ ]: class MulticlassPerceptron:
    """
    Implementation of Multiclass Perceptron using one-vs-rest strategy
    """

    def __init__(self):
        self.__weights = None
        self.__n_classes = None

    def __update(self, x, y, y_hat):
        """

```

```

Update weights for misclassified examples
s
Arg(s):
    x : numpy.ndarray
        Feature vector of shape  $d \times 1$ 
    y : int
        Label/target (-1 or 1)
    y_hat : int
        Predicted label (-1 or 1)
    ...

# DONE: Implement weight update rule for multiclass case
self.__weights[:, y] += x
self.__weights[:, y_hat] -= x

def fit(self, x, y, max_iter=100):
    """
    Fit the multiclass perceptron to training data

    Arg(s):
        x : numpy.ndarray
            Feature vector of shape  $d \times N$ 
        y : numpy.ndarray
            Label/target (-1 or 1) of shape  $(1 \times N)$  or  $(N \times 1)$ 
        max_iter : int
            Maximum number of iterations
    """

    # Flattening y for 1D everywhere
    y = y.flatten()

    n_features, n_samples = x.shape

    # DONE: Get number of classes from unique values in y
    self.__n_classes = np.unique(y).size

    # DONE: Initialize weights matrix of zeros with shape  $d+1 \times C$ 
    self.__weights = np.zeros((n_features + 1, self.__n_classes))

    # DONE: Append artificial coordinate ( $x_0$ ) to the data such that it is  $d+1 \times N$ 
    x = np.concatenate((np.ones((1, n_samples))), x), axis=0)

    # DONE: Implement training loop
    for _ in range(max_iter):
        n_updates = 0

```

```

        # Process each sample
        for n in range(n_samples):

            # DONE: Calculate scores and make prediction for each class
            scores = self.__weights.T @ x[:, n]
            y_hat = np.argmax(scores)

            # Update if prediction is wrong
            if y_hat != y[n]:
                self.__update(x[:, n], y[n], y_hat)
                n_updates += 1

            # DONE: Break if no updates were made, e.g., check for convergence
            if n_updates == 0:
                break

def predict(self, x):
    """
    Make predictions on new data

    Arg(s):
        x : numpy.ndarray
            Features of shape d x N

    Returns:
        numpy.ndarray : Predicted class labels
    """

    n_features, n_samples = x.shape

    # DONE: Append artificial coordinate (x0) to the data
    x = np.concatenate((np.ones((1, n_samples))), x), axis=0)

    # DONE: Implement prediction logic for multiclass case
    scores = self.__weights.T @ x
    multi_prediction = np.argmax(scores, axis=0).flatten()

    return multi_prediction

def score(self, x, y):
    """
    Calculate prediction accuracy

    Arg(s):
        x : numpy.ndarray
            Features of shape d x N
        y : numpy.ndarray

```

```

        Label/target (-1 or 1) of shape (1 x N) or (N x 1)

    Returns:
        float : Accuracy score
    """

    # DONE: Implement accuracy calculation
    multi_predictions = self.predict(x)
    multi_accuracy = np.mean(multi_predictions == y.flatten()) # Compare
    ↪with flattened y

    return multi_accuracy

```

```

[29]: def prepare_binary_digits_data(digits_zero=0, digits_one=1):
    """
    Prepare binary classification dataset from digits

    Args:
        digits_zero : int
            First digit to classify
        digits_one : int
            Second digit to classify

    Returns:
        tuple: (X_train, y_train, X_val, y_val, X_test, y_test)
            X_train : N x d
            y_train : N x 1
            X_val : M x d
            y_val : M x 1
            X_test : P x d
            y_test : P x 1
    """

    # Load digits dataset using sklearn.datasets
    digits = skdata.load_digits()

    # Select only the two specified digits
    mask = np.isin(digits.target, [digits_zero, digits_one])
    X = digits.data[mask]
    y = digits.target[mask]

    # Convert labels to -1/1
    y = np.where(y == digits_zero, -1, 1)

    # Split into train (60%), validation (20%), and test (20%) sets using
    ↪random_state=42

```

```

X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp,
↳test_size=0.25, random_state=42)

# Had to change a little to get it to work with the rest of the code (X.T
↳instead of X)
return X_train.T, np.expand_dims(y_train, axis=-1), X_val.T, np.
↳expand_dims(y_val, axis=-1), X_test.T, np.expand_dims(y_test, axis=-1)

```

```

[30]: def prepare_multiclass_digits_data():
    """
    Prepare multiclass classification dataset from digits

    Returns:
        tuple: (X_train, y_train, X_val, y_val, X_test, y_test)
            X_train : N x d
            y_train : N x 1
            X_val : M x d
            y_val : M x 1
            X_test : P x d
            y_test : P x 1
    """

    # Load digits dataset using sklearn.datasets
    digits = skdata.load_digits()
    X, y = digits.data, digits.target

    # Split into train (60%), validation (20%), and test (20%) sets with
    ↳random_state=42
    X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2,
    ↳random_state=42)
    X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp,
    ↳test_size=0.25, random_state=42)

    # NOTE: Had to change a little to get it to work with the rest of the code
    ↳ (X.T instead of X so shape is (d, N))
    return X_train.T, np.expand_dims(y_train, axis=-1), X_val.T, np.
    ↳expand_dims(y_val, axis=-1), X_test.T, np.expand_dims(y_test, axis=-1)

```

```

[ ]: # Binary classification experiment
print("Binary Classification Experiment (0 vs 1)")
print("-" * 50)

labels = [0, 1]

```



```

# Load and prepare binary data (0 vs 1)
X_train, y_train, X_val, y_val, X_test, y_test = prepare_binary_digits_data(0,
↪1)

# Try different max_iter values
max_iters = [10, 50, 100]
best_val_score = 0
best_model = None

for max_iter in max_iters:

    # DONE: Initialize and train binary perceptron
    model = BinaryPerceptron()
    model.fit(X_train, y_train)

    # DONE: Calculate validation score
    val_score = model.score(X_val, y_val)
    print("Max iterations: {}, Validation accuracy: {:.4f}".format(max_iter,
↪val_score))

    # DONE: Update best_model if current model performs better
    if val_score > best_val_score:
        best_val_score = val_score
        best_model = model

# DONE: Test best model on test set
test_score = best_model.score(X_test, y_test)
print("\nBest model test accuracy: {:.4f}".format(test_score))

# DONE: Create a confusion matrix using skmetrics.confusion_matrix on the test
↪set
# NOTE: Flattening since y_test is (P, 1)
binary_conf_matrix = skmetrics.confusion_matrix(y_test.flatten(), best_model.
↪predict(X_test))

# Show confusion matrix
binary_conf_matrix_plot = skmetrics.ConfusionMatrixDisplay(
    confusion_matrix=binary_conf_matrix,
    display_labels=labels
)
binary_conf_matrix_plot.plot()
plt.show()
time.sleep(1)

# DONE: Compare with scikit-learn implementation by training with max_iter=10
↪and random_state=42 and testing on the test set

```

```

# NOTE: Had to transpose X_train and X_test / flatten y_train and y_test to get
↳ the shapes to match to match sk docs
sk_model = Perceptron(max_iter=10, random_state=42)
sk_model.fit(X_train.T, y_train.flatten())
sk_score = sk_model.score(X_test.T, y_test.flatten())
print("Scikit-learn Perceptron test accuracy: {:.4f}".format(sk_score))

# DONE: Create a confusion matrix using skmetrics.confusion_matrix for scikit
↳ model on the test set
sk_binary_conf_matrix = skmetrics.confusion_matrix(y_test.flatten(), sk_model.
↳ predict(X_test.T))
sk_binary_conf_matrix_plot = skmetrics.ConfusionMatrixDisplay(
    confusion_matrix=sk_binary_conf_matrix,
    display_labels=labels
)
sk_binary_conf_matrix_plot.plot()
plt.show()
time.sleep(1)

print("\nMulticlass Classification Experiment (0-9)")
print("-" * 50)

labels = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

# Load and prepare multiclass data
X_train, y_train, X_val, y_val, X_test, y_test =
↳ prepare_multiclass_digits_data()

# Try different max_iter values
max_iters = [10, 50, 100]
best_val_score = 0
best_model = None

for max_iter in max_iters:

    # DONE: Initialize and train multiclass perceptron
    model = MulticlassPerceptron()
    model.fit(X_train, y_train)

    # DONE: Calculate validation score
    val_score = model.score(X_val, y_val)
    print("Max iterations: {}, Validation accuracy: {:.4f}".format(max_iter,
↳ val_score))

    # DONE: Update best_model if current model performs better
    if val_score > best_val_score:

```

```

        best_val_score = val_score
        best_model = model

# DONE: Test best model on test set
test_score = best_model.score(X_test, y_test)
print("\nBest model test accuracy: {:.4f}".format(test_score))

# DONE: Create a confusion matrix using skmetrics.confusion_matrix for your
↳model on the test set
multi_conf_matrix = skmetrics.confusion_matrix(y_test.flatten(), best_model.
↳predict(X_test))
multi_conf_matrix_plot = skmetrics.ConfusionMatrixDisplay(
    confusion_matrix=multi_conf_matrix,
    display_labels=labels
)
multi_conf_matrix_plot.plot()
plt.show()
time.sleep(1)

# DONE: Compare with scikit-learn implementation by training with max_iter=10
↳and random_state=42 and testing on the test set
# NOTE: Had to transpose X_test and X_train / flatten y_train and y_test to get
↳the shapes to match to match sk docs
sk_model = Perceptron(max_iter=10, random_state=42)

sk_model.fit(X_train.T, y_train.flatten())
sk_score = sk_model.score(X_test.T, y_test.flatten())
print("Scikit-learn Perceptron test accuracy: {:.4f}".format(sk_score))

# DONE: Create a confusion matrix using skmetrics.confusion_matrix for your
↳model on the test set
sk_multi_conf_matrix = skmetrics.confusion_matrix(y_test.flatten(), sk_model.
↳predict(X_test.T))

# Show confusion matrix
sk_multi_conf_matrix_plot = skmetrics.ConfusionMatrixDisplay(
    confusion_matrix=sk_multi_conf_matrix,
    display_labels=labels
)
sk_multi_conf_matrix_plot.plot()
plt.show()
time.sleep(1)

```


Binary Classification Experiment (0 vs 1)

Max iterations: 10, Validation accuracy: 1.0000


Max iterations: 50, Validation accuracy: 1.0000

Max iterations: 100, Validation accuracy: 1.0000

Best model test accuracy: 0.9861



Scikit-learn Perceptron test accuracy: 1.0000



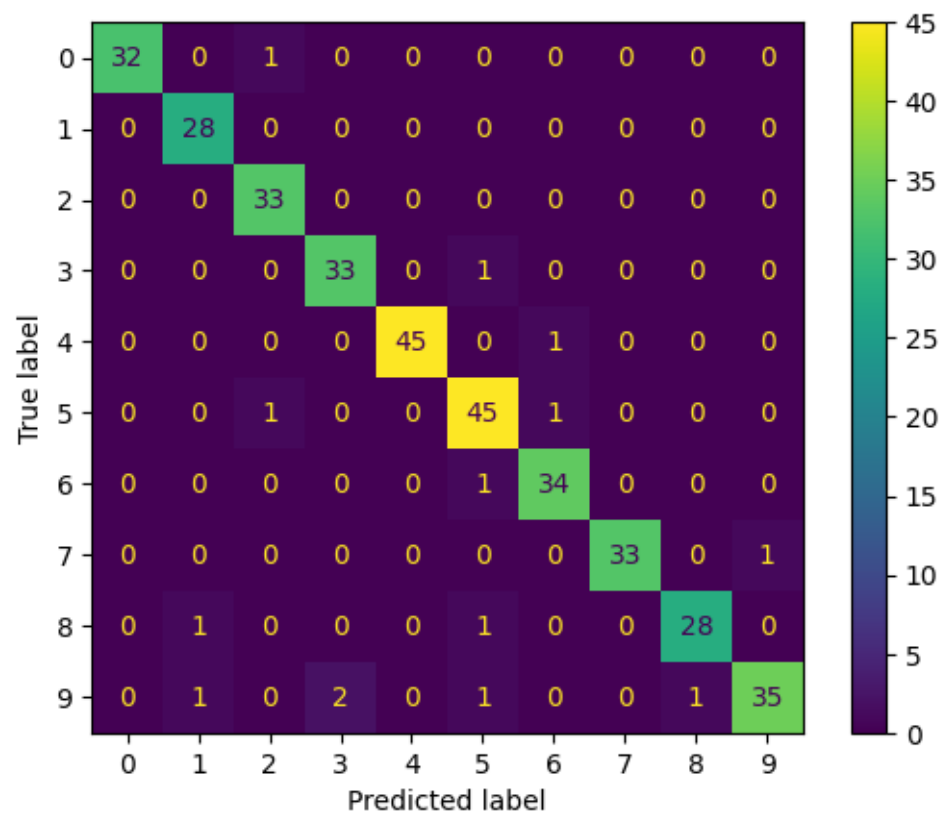
Multiclass Classification Experiment (0-9)

Max iterations: 10, Validation accuracy: 0.9583

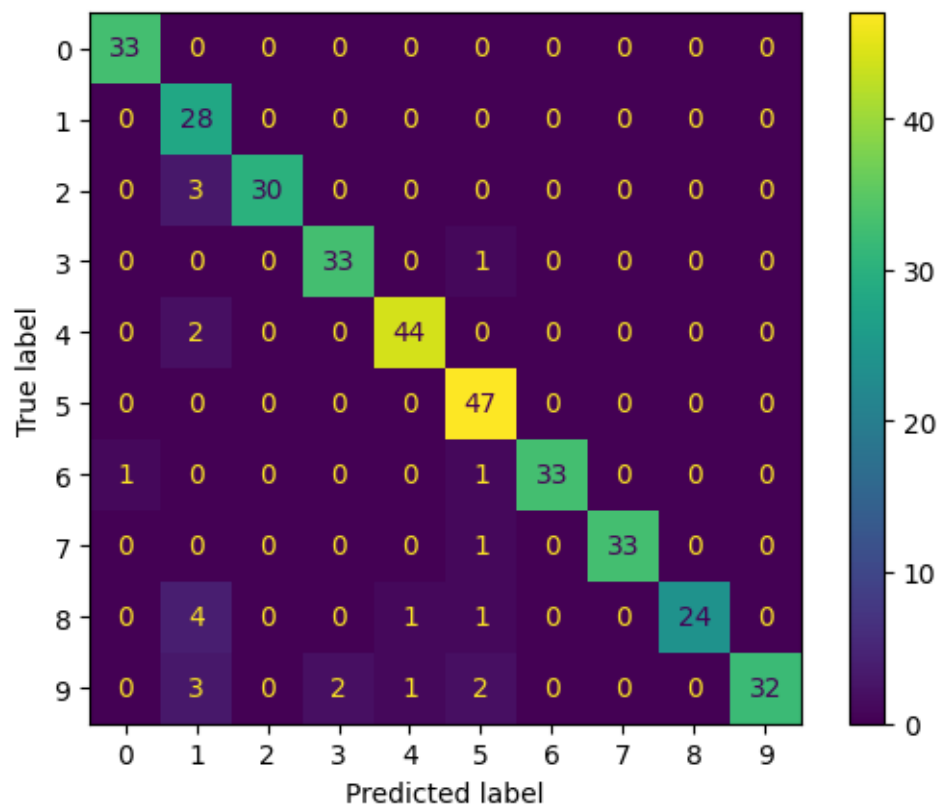
Max iterations: 50, Validation accuracy: 0.9583

Max iterations: 100, Validation accuracy: 0.9583

Best model test accuracy: 0.9611



Scikit-learn Perceptron test accuracy: 0.9361



[]:

[]: