

02_assignment_stochastic_gradient_descent_robertson1

March 24, 2025

Assignment 2: Stochastic Gradient Descent and Momentum

CPSC 381/581: Machine Learning

Yale University

Instructor: Alex Wong

Student: Hailey Robertson

Prerequisites:

1. Enable Google Colaboratory as an app on your Google Drive account
2. Create a new Google Colab notebook, this will also create a “Colab Notebooks” directory under “MyDrive” i.e.

`/content/drive/MyDrive/Colab Notebooks`

3. Create the following directory structure in your Google Drive

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments`

4. Move the `02_assignment_stochastic_gradient_descent.ipynb` into

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments`

so that its absolute path is

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments/02_assignment`

In this assignment, we will optimize a linear function for the logistic regression task using the stochastic gradient descent and its momentum variant. We will test them on several binary classification datasets (breast cancer, digits larger or less than 5, and fir and pine coverage). We will implement a training and validation loop for the binary classification task and test it on the testing split for each dataset.

Submission:

1. Implement all TODOs in the code blocks below.
2. Report your training, validation, and testing scores.

Report training, validation, and testing scores here.

3. List any collaborators.

Collaborators: N/A

IMPORTANT:

- For full credit, your mean classification accuracies for all trained models across all datasets should be no more than 8% worse the scores achieved by sci-kit learn's logistic regression model across training, validation and testing splits.
- You may not use batch sizes of more than 10% of the dataset size for stochastic gradient descent and momentum stochastic gradient descent.
- You will only need to experiment with gradient descent (GD) and momentum gradient descent (momentum GD) on breast cancer and digits (toy) datasets. It will take too long to run them on fir and pine coverage (realistic) dataset to get reasonable numbers. Of course, you may try them on fir and pine coverage :) but they will not count towards your grade.
- Note the run time speed up when comparing GD and momentum GD with stochastic gradient descent (SGD) and momentum stochastic gradient descent (momentum SGD)! Even though they are faster and observing batches instead of the full dataset at each time step, they can still achieving similar accuracies!

```
[6]: import numpy as np
import sklearn.datasets as skdata
import sklearn.metrics as skmetrics
from sklearn.linear_model import LogisticRegression as LogisticRegressionSciKit
import warnings
import time

warnings.filterwarnings(action='ignore')
np.random.seed = 1
```

Implementation of stochastic gradient descent optimizer for logistic loss

```
[7]: class Optimizer(object):

    def __init__(self, alpha, eta_decay_factor, beta, optimizer_type):
        '''
        Arg(s):
            alpha : float
                    initial learning rate
            eta_decay_factor : float
                    learning rate decay rate
            beta : float
                    momentum discount rate
            optimizer_type : str
                    'gradient_descent',
                    'momentum_gradient_descent',
                    'stochastic_gradient_descent',
                    'momentum_stochastic_gradient_descent'
        '''

        self.__alpha = alpha
```

```

self.__eta_decay_factor = eta_decay_factor
self.__beta = beta
self.__optimizer_type = optimizer_type
self.__momentum = None

def __compute_gradients(self, w, x, y, loss_func='logistic'):
    """
    Returns the gradient of a loss function

    Arg(s):
        w : numpy[float32]
            d x 1 weight vector
        x : numpy[float32]
            d x N feature vector
        y : numpy[float32]
            1 x N groundtruth vector
        loss_func : str
            loss by default is 'logistic' only for the purpose of the
    """
    assignment
    Returns:
        numpy[float32] : d x 1 gradients
    """

    # DONE: Implement compute_gradient function

    if loss_func == 'logistic':
        # print('w shape: {}'.format(w.shape)) # (1, 30)
        N = x.shape[1] # number of samples, 341
        z = np.dot(w.T, x)
        # print('y shape: {}'.format(y.shape)) # (1, 341)
        # print('z shape: {}'.format(z.shape)) # (1, 341)
        # sigmoid = 1 / (1 + np.exp(-y * z))
        # gradient = -np.dot(x, ((y * (1 - sigmoid))).T) / N
        gradient = -np.mean((y * x) / (1 + np.exp(y * z)), axis=1,
    keepdims=True)

        # print('gradient shape: {}'.format(gradient.shape))

        return gradient

    else:
        raise ValueError('Unsupported loss function: {}'.format(loss_func))

def __polynomial_decay(self, time_step):
    """
    Computes the polynomial decay factor  $t^{-a}$ 

```

```

    Arg(s):
        time_step : int
            current step in optimization
    Returns:
        float : polynomial decay to adjust (reduce) initial learning rate
        ...

# DONE: Implement polynomial decay to adjust the initial learning rate
if self.__eta_decay_factor is None:
    return self.__alpha
else:
    decay_factor = time_step ** (-self.__eta_decay_factor)
    decayed_alpha = self.__alpha * decay_factor

    return decayed_alpha

def update(self,
           w,
           x,
           y,
           loss_func,
           batch_size,
           time_step):
    """
    Updates the weight vector based on

    Arg(s):
        w : numpy[float32]
            d x 1 weight vector
        x : numpy[float32]
            d x N feature vector
        y : numpy[float32]
            1 x N groundtruth vector
        loss_func : str
            loss function to use, should be 'logistic' for the purpose of
↳ the assignment
        batch_size : int
            batch size for stochastic and momentum stochastic gradient
↳ descent
        time_step : int
            current step in optimization
    Returns:
        numpy[float32]: d x 1 weights
        ...

# DONE: Implement the optimizer update function

```

```

# For each optimizer type, compute gradients and update weights
eta = self.__polynomial_decay(time_step) # step size
gradient = self.__compute_gradients(w, x, y, loss_func)

if self.__optimizer_type == 'gradient_descent':
    # print('gradient shape: {}'.format(gradient.shape)) # (30, 1)
    # print('w shape: {}'.format(w.shape)) # (30, 1)
    # print("Before update: w[0]:", w[0])
    w = w - eta * gradient
    # print("After update: w[0]:", w[0])
    return w

elif self.__optimizer_type == 'momentum_gradient_descent':
    if self.__momentum is None:
        self.__momentum = np.zeros(w.shape)
    self.__momentum = self.__beta * self.__momentum + (1-self.__beta) * ↵
    ↵gradient
    w = w - eta * self.__momentum
    return w

elif self.__optimizer_type == 'stochastic_gradient_descent':
    indices = np.random.choice(x.shape[1], batch_size, replace=False)
    x_batch = x[:, indices]
    y_batch = y[:, indices]
    gradient = self.__compute_gradients(w, x_batch, y_batch, loss_func)
    w = w - eta * gradient
    return w

elif self.__optimizer_type == 'momentum_stochastic_gradient_descent':
    if self.__momentum is None:
        self.__momentum = np.zeros(w.shape)
    indices = np.random.choice(x.shape[1], batch_size, replace=False)
    x_batch = x[:, indices]
    y_batch = y[:, indices]
    gradient = self.__compute_gradients(w, x_batch, y_batch, loss_func)
    self.__momentum = self.__beta * self.__momentum + (1-self.__beta) * ↵
    ↵gradient
    w = w - eta * self.__momentum
    return w

else:
    raise ValueError('Unsupported optimizer type: {}'.format(self.
    ↵__optimizer_type))

```

Implementation of our logistic regression model for binary classification

```
[8]: class LogisticRegression(object):

    def __init__(self):
        # Define private variables
        self.__weights = None
        self.__optimizer = None

    def fit(self,
            x,
            y,
            T,
            alpha,
            eta_decay_factor,
            beta,
            batch_size,
            optimizer_type,
            loss_func='logistic'):
        '''
        Fits the model to x and y by updating the weight vector
        using gradient descent

        Arg(s):
            x : numpy[float32]
                d x N feature vector
            y : numpy[float32]
                1 x N groundtruth vector
            T : int
                number of iterations to train
            alpha : float
                learning rate
            eta_decay_factor : float
                learning rate decay rate
            beta : float
                momentum discount rate
            batch_size : int
                number of examples per batch
            optimizer_type : str
                'gradient_descent',
                'momentum_gradient_descent',
                'stochastic_gradient_descent',
                'momentum_stochastic_gradient_descent'
            loss_func : str
                loss function to use, by default is 'logistic' only for the
                purpose of the assignment
        '''

        # DONE: Instantiate optimizer and weights
```

```

        self.__optimizer = Optimizer(alpha, eta_decay_factor, beta,
↪optimizer_type)
        self.__weights = np.random.randn(x.shape[0], 1) * 0.01

    for t in range(1, T + 1):

        # DONE: Compute loss function
        loss = self.__compute_loss(x, y, loss_func)

        if (t % 1000) == 0:
            print('Step={} Loss={}'.format(t, loss))

        # DONE: Update weights
        self.__weights = self.__optimizer.update(self.__weights, x, y,
↪loss_func, batch_size, t)
        # print("loss", loss, "weights", self.__weights)

    def predict(self, x):
        """
        Predicts the label for each feature vector x

        Arg(s):
            x : numpy[float32]
                d x N feature vector
        Returns:
            numpy[float32] : 1 x N vector
        """

        # DONE: Implements the predict function
        # Hint: logistic regression predicts a value between 0 and 1
        z = np.dot(self.__weights.T, x)
        sigmoid = 1 / (1 + np.exp(-z))
        # print(sigmoid.shape)
        predictions = np.where(sigmoid >= 0.5, 1, -1)
        # print(predictions)

        return predictions

    def __compute_loss(self, x, y, loss_func):
        """
        Computes the logistic loss

        Arg(s):
            x : numpy[float32]
                d x N feature vector

```

```

        y : numpy[float32]
            1 x N groundtruth vector
        loss_func : str
            loss function to use, by default is 'logistic' only for the
        ↪ purpose of the assignment
    Returns:
        float : loss
    '''

    # DONE: Implements the __compute_loss function
    if loss_func == 'logistic':

        N = x.shape[1] # number of samples, 341
        z = np.dot(self.__weights.T, x)
        loss = np.sum(np.log(1 + np.exp(-y * z))) / N

    else:
        raise ValueError('Unsupported loss function: {}'.format(loss_func))

    return loss

```

Training, validating and testing logistic regression for binary classification

```

[9]: def compare_scores(score_train, score_val, score_test, sk_score_train,
    ↪ sk_score_val, sk_score_test):
    """
        Compare the performance of our logistic regression model with scikit-learn
    ↪ model.
        If the performance is more than 8% worse for any dataset (train/validation/
    ↪ test),
        print 'Rerun' because I don't want to assess this by hand.
    """
    threshold = 0.08

    if score_train < sk_score_train * (1 - threshold):
        print(f"\nTraining score is more than 8% worse. Ours: {score_train:.
    ↪ 4f}, Scikit: {sk_score_train:.4f}. Rerun.")

    if score_val < sk_score_val * (1 - threshold):
        print(f"\nValidation score is more than 8% worse. Ours: {score_val:.
    ↪ 4f}, Scikit: {sk_score_val:.4f}. Rerun.")

    if score_test < sk_score_test * (1 - threshold):
        print(f"\nTest score is more than 8% worse. Ours: {score_test:.4f},
    ↪ Scikit: {sk_score_test:.4f}. Rerun.")

    if (score_train >= sk_score_train * (1 - threshold)) and \

```



```

(score_val >= sk_score_val * (1 - threshold)) and \
(score_test >= sk_score_test * (1 - threshold)):
    print("\nPerformance is acceptable.\n")

```

```

[10]: # Load breast cancer, digits, and tree coverage datasets
datasets = [
    skdata.load_breast_cancer(),
    skdata.load_digits(),
    skdata.fetch_covtype()
]
dataset_names = [
    'breast cancer',
    'digits greater or less than 5',
    'fir and pine coverage',
]

# Loss functions to minimize
dataset_optimizer_types = [
    # For breast cancer dataset
    [
        'gradient_descent',
        'momentum_gradient_descent',
        'stochastic_gradient_descent',
        'momentum_stochastic_gradient_descent'
    ],
    # For digits greater than or less than 5 dataset
    [
        'gradient_descent',
        'momentum_gradient_descent',
        'stochastic_gradient_descent',
        'momentum_stochastic_gradient_descent'
    ],
    # For fir and pine coverage dataset
    [
        'stochastic_gradient_descent',
        'momentum_stochastic_gradient_descent'
    ]
]

# DONE: Select hyperparameters

# Step size (always used)
dataset_alphas = [
    # For breast cancer dataset
    [1e-6, 1e-6, 0.001, 0.01],
    # For digits greater than or less than 5 dataset
    [1e-4, 1e-4, 0.01, 0.1],
    # For fir and pine coverage dataset
    [0.001, 0.01]
]

```

```

]

# How much the learning rate should decrease over time (only for SGD)
dataset_eta_decay_factors = [
    # For breast cancer dataset
    [None, None, 0.99, 0.99],
    # For digits greater than or less than 5 dataset
    [None, None, 0.99, 0.99],
    # For fir and pine coverage dataset
    [0.99, 0.99]
]

# Momentum (only for momentum-based ..duh)
# None for no momentum
dataset_betas = [
    # For breast cancer dataset
    [None, 0.9, None, 0.9],
    # For digits greater than or less than 5 dataset
    [None, 0.9, None, 0.9],
    # For fir and pine coverage dataset
    [None, 0.9]
]

# Samples (only for SGD)
dataset_batch_sizes = [
    # For breast cancer dataset
    # # N = 569
    [None, None, 24, 24],
    # For digits greater than or less than 5 dataset
    # # N = 1797
    [None, None, 32, 32],
    # For fir and pine coverage dataset
    # # N = 495141
    [1000, 1000]
]

# Iterations
dataset_Ts = [
    # For breast cancer dataset
    [10000, 10000, 10000, 10000],
    # For digits greater than or less than 5 dataset
    [10000, 10000, 10000, 10000],
    # For fir and pine coverage dataset
    [3000, 3000]
]

```

```

# Zip up all dataset options
dataset_options = zip(
    datasets,
    dataset_names,
    dataset_optimizer_types,
    dataset_alphas,
    dataset_eta_decay_factors,
    dataset_betas,
    dataset_batch_sizes,
    dataset_Ts)

for options in dataset_options:

    # Unpack dataset options
    dataset, \
        dataset_name, \
        optimizer_types, \
        alphas, \
        eta_decay_factors, \
        betas, \
        batch_sizes, \
        Ts = options

    '''
    Create the training, validation and testing splits
    '''

    x = dataset.data
    y = dataset.target

    if dataset_name == 'digits greater or less than 5':
        y[y < 5] = 1
        y[y >= 5] = 0
    elif dataset_name == 'fir and pine coverage':

        idx_fir_or_pine = np.where(np.logical_or(y == 1, y == 2))[0]

        x = x[idx_fir_or_pine, :]
        y = y[idx_fir_or_pine]

        # Pine class: 0; Fir class: 1
        y[y == 2] = 0

    print('Preprocessing the {} dataset ({} samples, {} feature dimensions)'.
    ↪format(dataset_name, x.shape[0], x.shape[1]))

    # Shuffle the dataset based on sample indices

```

```

shuffled_indices = np.random.permutation(x.shape[0])

# Choose the first 60% as training set, next 20% as validation and the rest
↳as testing
train_split_idx = int(0.60 * x.shape[0])
val_split_idx = int(0.80 * x.shape[0])

train_indices = shuffled_indices[0:train_split_idx]
val_indices = shuffled_indices[train_split_idx:val_split_idx]
test_indices = shuffled_indices[val_split_idx:]

# Select the examples from x and y to construct our training, validation,
↳testing sets
x_train, y_train = x[train_indices, :], y[train_indices]
x_val, y_val = x[val_indices, :], y[val_indices]
x_test, y_test = x[test_indices, :], y[test_indices]

'''
Trains and tests logistic regression model from scikit-learn
'''
model_scikit = LogisticRegressionSciKit(penalty=None, fit_intercept=False)

# DONE: Train scikit-learn logistic regression model
model_scikit.fit(x_train, y_train)

print('***** Results on the {} dataset using scikit-learn logistic
↳regression model *****'.format(dataset_name))

# DONE: Score model using mean accuracy on training set
predictions_train = model_scikit.predict(x_train)
sk_score_train = skmetrics.accuracy_score(y_train, predictions_train)
print('Training set mean accuracy: {:.4f}'.format(sk_score_train))

# DONE: Score model using mean accuracy on validation set
predictions_val = model_scikit.predict(x_val)
sk_score_val = skmetrics.accuracy_score(y_val, predictions_val)
print('Validation set mean accuracy: {:.4f}'.format(sk_score_val))

# DONE: Score model using mean accuracy on testing set
predictions_test = model_scikit.predict(x_test)
sk_score_test = skmetrics.accuracy_score(y_test, predictions_test)
print('Testing set mean accuracy: {:.4f}'.format(sk_score_test))

'''
Trains, validates, and tests our logistic regression model for binary
↳classification

```

```

'''
    # Take the transpose of the dataset to match the dimensions discussed in
    ↳lecture
    # i.e., (N x d) to (d x N)
    x_train = np.transpose(x_train, axes=(1, 0))
    x_val = np.transpose(x_val, axes=(1, 0))
    x_test = np.transpose(x_test, axes=(1, 0))
    y_train = np.expand_dims(y_train, axis=0)
    y_val = np.expand_dims(y_val, axis=0)
    y_test = np.expand_dims(y_test, axis=0)

    # DONE: Set the ground truth to the appropriate classes (integers)
    ↳according to lecture
    # -1 and +1
    y_train = np.where(y_train == 0, -1, 1)
    y_val = np.where(y_val == 0, -1, 1)
    y_test = np.where(y_test == 0, -1, 1)

    model_options = zip(optimizer_types, alphas, eta_decay_factors, betas,
    ↳batch_sizes, Ts)

    for optimizer_type, alpha, eta_decay_factor, beta, batch_size, T in
    ↳model_options:

        # DONE: Initialize our logistic regression model
        model_ours = LogisticRegression()

        print('***** Results of our logistic regression model trained on {}
        ↳dataset *****'.format(dataset_name))
        print('\t optimizer_type={} \n\t alpha={} \n\t eta_decay_factor={} \n\t
        ↳beta={} \n\t batch_size={} \n\t T={} '.format(
            optimizer_type, alpha, eta_decay_factor, beta, batch_size, T))

        time_start = time.time()

        # DONE: Train model on training set
        model_ours.fit(
            x_train,
            y_train,
            T,
            alpha,
            eta_decay_factor,
            beta,
            batch_size,
            optimizer_type
        )

```

```

time_elapsed = time.time() - time_start
print('Total training time: {:3f} seconds'.format(time_elapsed))

# DONE: Score model using mean accuracy on training set
predictions_train = model_ours.predict(x_train)
score_train = np.mean(predictions_train == y_train)
# print("Predictions", predictions_train)
# print("Ground truth", y_train)
print('Training set mean accuracy: {:.4f}'.format(score_train))

# DONE: Score model using mean accuracy on training set
predictions_val = model_ours.predict(x_val)
score_val = np.mean(predictions_val == y_val)
print('Validation set mean accuracy: {:.4f}'.format(score_val))

# DONE: Score model using mean accuracy on training set
predictions_test = model_ours.predict(x_test)
score_test = np.mean(predictions_test == y_test)
print('Testing set mean accuracy: {:.4f}'.format(score_test))

compare_scores(
    score_train, score_val, score_test,
    sk_score_train, sk_score_val, sk_score_test
)

print('')

```

Preprocessing the breast cancer dataset (569 samples, 30 feature dimensions)
 ***** Results on the breast cancer dataset using scikit-learn logistic regression model *****
 Training set mean accuracy: 0.9413
 Validation set mean accuracy: 0.9649
 Testing set mean accuracy: 0.9211
 ***** Results of our logistic regression model trained on breast cancer dataset *****

```

optimizer_type=gradient_descent
alpha=1e-06
eta_decay_factor=None
beta=None
batch_size=None
T=10000
Step=1000 Loss=0.46044808320071234
Step=2000 Loss=0.3804614825325583
Step=3000 Loss=0.3407873696823556
Step=4000 Loss=0.31573874221992476
Step=5000 Loss=0.29815859234079
Step=6000 Loss=0.2850233088381083

```

```
Step=7000   Loss=0.2747774657667066
Step=8000   Loss=0.26652762410995945
Step=9000   Loss=0.2597214198332295
Step=10000  Loss=0.2539980216090421
Total training time: 0.475812 seconds
Training set mean accuracy: 0.9208
Validation set mean accuracy: 0.9386
Testing set mean accuracy: 0.9035
```

Performance is acceptable.

***** Results of our logistic regression model trained on breast cancer dataset

```
optimizer_type=momentum_gradient_descent
alpha=1e-06
eta_decay_factor=None
beta=0.9
batch_size=None
T=10000
```

```
Step=1000   Loss=0.46930506241203157
Step=2000   Loss=0.4049297795647417
Step=3000   Loss=0.3606284108999625
Step=4000   Loss=0.32999529856826276
Step=5000   Loss=0.30884989557851533
Step=6000   Loss=0.2934165449902679
Step=7000   Loss=0.2814993346181599
Step=8000   Loss=0.2719409595126229
Step=9000   Loss=0.2640711888086081
Step=10000  Loss=0.25746568486879196
Total training time: 0.487436 seconds
Training set mean accuracy: 0.9150
Validation set mean accuracy: 0.9386
Testing set mean accuracy: 0.9035
```

Performance is acceptable.

***** Results of our logistic regression model trained on breast cancer dataset

```
optimizer_type=stochastic_gradient_descent
alpha=0.001
eta_decay_factor=0.99
beta=None
batch_size=24
T=10000
```

```
Step=1000   Loss=0.24863761232744666
Step=2000   Loss=0.23934020827658284
Step=3000   Loss=0.23595073660498928
Step=4000   Loss=0.23336870679176538
```

```
Step=5000   Loss=0.23172014163508645
Step=6000   Loss=0.23052145859595313
Step=7000   Loss=0.22958090429883213
Step=8000   Loss=0.22877424878861352
Step=9000   Loss=0.22826028798798242
Step=10000  Loss=0.22762238137516555
Total training time: 0.714931 seconds
Training set mean accuracy: 0.9091
Validation set mean accuracy: 0.9561
Testing set mean accuracy: 0.8947
```

Performance is acceptable.

***** Results of our logistic regression model trained on breast cancer dataset

```
optimizer_type=momentum_stochastic_gradient_descent
alpha=0.01
eta_decay_factor=0.99
beta=0.9
batch_size=24
T=10000
Step=1000   Loss=0.9473641861922635
Step=2000   Loss=0.88717325106249
Step=3000   Loss=0.793637495810315
Step=4000   Loss=0.7609227853764511
Step=5000   Loss=0.7435816907194572
Step=6000   Loss=0.7207641303526582
Step=7000   Loss=0.7060942422441724
Step=8000   Loss=0.6960712284775495
Step=9000   Loss=0.6829020556972859
Step=10000  Loss=0.6756001298112138
Total training time: 0.728675 seconds
Training set mean accuracy: 0.9032
Validation set mean accuracy: 0.8947
Testing set mean accuracy: 0.8684
```

Performance is acceptable.

Preprocessing the digits greater or less than 5 dataset (1797 samples, 64 feature dimensions)

***** Results on the digits greater or less than 5 dataset using scikit-learn logistic regression model *****

```
Training set mean accuracy: 0.9239
Validation set mean accuracy: 0.8552
Testing set mean accuracy: 0.8917
```

***** Results of our logistic regression model trained on digits greater or less than 5 dataset *****


```

optimizer_type=gradient_descent
alpha=0.0001
eta_decay_factor=None
beta=None
batch_size=None
T=10000
Step=1000 Loss=0.4129632392476915
Step=2000 Loss=0.34781596948087756
Step=3000 Loss=0.3167059121793643
Step=4000 Loss=0.29754526750808924
Step=5000 Loss=0.28422129329587076
Step=6000 Loss=0.2742877485846041
Step=7000 Loss=0.2665430919465577
Step=8000 Loss=0.260314388924469
Step=9000 Loss=0.25518885078340586
Step=10000 Loss=0.2508959073545494
Total training time: 1.614091 seconds
Training set mean accuracy: 0.9091
Validation set mean accuracy: 0.8468
Testing set mean accuracy: 0.9056

```

Performance is acceptable.

***** Results of our logistic regression model trained on digits greater or less than 5 dataset *****

```

optimizer_type=momentum_gradient_descent
alpha=0.0001
eta_decay_factor=None
beta=0.9
batch_size=None
T=10000
Step=1000 Loss=0.4121279174215425
Step=2000 Loss=0.34589851792735093
Step=3000 Loss=0.3150581979577341
Step=4000 Loss=0.29623090338873154
Step=5000 Loss=0.2831840764287226
Step=6000 Loss=0.2734707065272934
Step=7000 Loss=0.2659012155295467
Step=8000 Loss=0.25981336833861285
Step=9000 Loss=0.2548024572413244
Step=10000 Loss=0.2506038501901871
Total training time: 1.605079 seconds
Training set mean accuracy: 0.9082
Validation set mean accuracy: 0.8496
Testing set mean accuracy: 0.9083

```

Performance is acceptable.

***** Results of our logistic regression model trained on digits greater or less than 5 dataset *****

optimizer_type=stochastic_gradient_descent
alpha=0.01
eta_decay_factor=0.99
beta=None
batch_size=32
T=10000

Step=1000 Loss=0.4391554060252024
Step=2000 Loss=0.42926102081513023
Step=3000 Loss=0.4238914588168506
Step=4000 Loss=0.4202584956032843
Step=5000 Loss=0.4175778652136804
Step=6000 Loss=0.4154192921630118
Step=7000 Loss=0.41365074551754644
Step=8000 Loss=0.41213248723079426
Step=9000 Loss=0.41082529289944525
Step=10000 Loss=0.40967846471837605
Total training time: 1.947454 seconds
Training set mean accuracy: 0.8627
Validation set mean accuracy: 0.8050
Testing set mean accuracy: 0.8806

Performance is acceptable.

***** Results of our logistic regression model trained on digits greater or less than 5 dataset *****

optimizer_type=momentum_stochastic_gradient_descent
alpha=0.1
eta_decay_factor=0.99
beta=0.9
batch_size=32
T=10000

Step=1000 Loss=0.2728942261913827
Step=2000 Loss=0.26714562266133673
Step=3000 Loss=0.2642149352245792
Step=4000 Loss=0.26227669352216093
Step=5000 Loss=0.2608649426573262
Step=6000 Loss=0.2597924350627642
Step=7000 Loss=0.2588988467798901
Step=8000 Loss=0.258137106510277
Step=9000 Loss=0.25753583901078986
Step=10000 Loss=0.2569099063125447
Total training time: 1.928692 seconds
Training set mean accuracy: 0.9100
Validation set mean accuracy: 0.8496
Testing set mean accuracy: 0.9028

Performance is acceptable.

Preprocessing the fir and pine coverage dataset (495141 samples, 54 feature dimensions)

***** Results on the fir and pine coverage dataset using scikit-learn logistic regression model *****

Training set mean accuracy: 0.7576

Validation set mean accuracy: 0.7606

Testing set mean accuracy: 0.7566

***** Results of our logistic regression model trained on fir and pine coverage dataset *****

optimizer_type=stochastic_gradient_descent

alpha=0.001

eta_decay_factor=0.99

beta=None

batch_size=1000

T=3000

Step=1000 Loss=1.9504758124222066

Step=2000 Loss=0.5990504055646176

Step=3000 Loss=0.5628246412609968

Total training time: 315.173391 seconds

Training set mean accuracy: 0.7096

Validation set mean accuracy: 0.7098

Testing set mean accuracy: 0.7068

Performance is acceptable.

***** Results of our logistic regression model trained on fir and pine coverage dataset *****

optimizer_type=momentum_stochastic_gradient_descent

alpha=0.01

eta_decay_factor=0.99

beta=0.9

batch_size=1000

T=3000

Step=1000 Loss=0.9902832334635768

Step=2000 Loss=0.8679242075417812

Step=3000 Loss=0.8040880543020845

Total training time: 304.039843 seconds

Training set mean accuracy: 0.7248

Validation set mean accuracy: 0.7249

Testing set mean accuracy: 0.7230

Performance is acceptable.

[]:

[]: