

04_exercise_gradient_descent_robertson

February 17, 2025

Exercise 4: Gradient Descent for Linear Regression

CPSC 381/581: Machine Learning

Yale University

Instructor: Alex Wong

Student: Hailey Robertson

Prerequisites:

1. Enable Google Colaboratory as an app on your Google Drive account
2. Create a new Google Colab notebook, this will also create a “Colab Notebooks” directory under “MyDrive” i.e.

`/content/drive/MyDrive/Colab Notebooks`

3. Create the following directory structure in your Google Drive

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Exercises`

4. Move the 04_exercise_gradient_descent.ipynb into

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Exercises`

so that its absolute path is

`/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Exercises/04_exercise_gradient_descent.ipynb`

In this exercise, we will optimize a linear function for the regression task using the gradient descent for mean squared and half mean squared losses. We will test them on several datasets.

Submission:

1. Implement all TODOs in the code blocks below.
2. Report your training, validation, and testing scores.

Report validation and testing scores here.

For full credit, your mean squared error scores for models optimized using mean_squared and half_mean_squared losses should be within 10% of the scores below.

***** Results of scikit-learn linear regression model on Diabetes dataset *****

Training set mean squared error: 25472.6663

Training set r-squared scores: -3.2907

Validation set mean squared error: 30376.7403

```

Validation set r-squared scores: -5.1176
Testing set mean squared error: 27024.0792
Testing set r-squared scores: -3.2604
***** Results of our linear regression model trained with mean_squared loss, alpha=0.0001 and T
Training set mean squared error: 28034.7946
Training set r-squared scores: -3.7223
Validation set mean squared error: 34538.4871
Validation set r-squared scores: -5.9557
Testing set mean squared error: 29713.4975
Testing set r-squared scores: -3.6844
***** Results of our linear regression model trained with half_mean_squared loss, alpha=0.0002
Training set mean squared error: 28034.7946
Training set r-squared scores: -3.7223
Validation set mean squared error: 34538.4871
Validation set r-squared scores: -5.9557
Testing set mean squared error: 29713.4975
Testing set r-squared scores: -3.6844
***** Results of scikit-learn linear regression model on California housing prices dataset ***
Training set mean squared error: 0.5989
Training set r-squared scores: 0.5463
Validation set mean squared error: 0.5922
Validation set r-squared scores: 0.5721
Testing set mean squared error: 0.6590
Testing set r-squared scores: 0.5185
***** Results of our linear regression model trained with mean_squared loss, alpha=1e-07 and T
Training set mean squared error: 0.6496
Training set r-squared scores: 0.5080
Validation set mean squared error: 0.6348
Validation set r-squared scores: 0.5413
Testing set mean squared error: 0.7028
Testing set r-squared scores: 0.4865
***** Results of our linear regression model trained with half_mean_squared loss, alpha=2e-07
Training set mean squared error: 0.6496
Training set r-squared scores: 0.5080
Validation set mean squared error: 0.6348
Validation set r-squared scores: 0.5413
Testing set mean squared error: 0.7028
Testing set r-squared scores: 0.4865

```

3. List any collaborators.

Collaborators: Doe, Jane (Please write names in <Last Name, First Name> format)

Collaboration details: Discussed ... implementation details with Jane Doe.

Import packages

```
[1]: import numpy as np
import sklearn.datasets as skdata
```

```

import sklearn.metrics as skmetrics
from sklearn.linear_model import LinearRegression as LinearRegressionSciKit
import warnings

warnings.filterwarnings(action='ignore')
np.random.seed = 1

```

Implementation of our Gradient Descent optimizer for mean squared and half mean squared loss

```

[2]: class GradientDescentOptimizer(object):

    def __init__(self):
        pass

    def __compute_gradients(self, w, x, y, loss_func):
        '''
        Returns the gradient of mean squared or half mean squared loss

        Arg(s):
            w : numpy[float32]
                d x 1 weight vector
            x : numpy[float32]
                d x N feature vector
            y : numpy[float32]
                1 x N groundtruth vector
            loss_func : str
                loss type either 'mean_squared', or 'half_mean_squared'
        Returns:
            numpy[float32] : d x 1 gradients
        '''

        # DONE: Implements the __compute_gradients function
        if loss_func == 'mean_squared':
            # DONE: Implements gradients for mean squared loss

            '''
            Using for-loop

            gradients = np.zeros(x.shape)

            for n in range((x.shape[1])):
                x_n = x[:, n]
                gradients[:, n] = (np.matmul(w.T, x_n) - y[n]) * x_n
            '''

            # Using matrix multiplication
            gradients = (np.matmul(w.T, x) - y) * x

```

```

        # Note: Set keepdims=True to keep the dimension of 1 (otherwise it
        ↪will get squashed by mean operation)
        return 2.0 * np.mean(gradients, axis=1, keepdims=True)
    elif loss_func == 'half_mean_squared':
        # DONE: Implements gradients for half mean squared loss

        gradients = (np.matmul(w.T, x) - y) * x

        return np.mean(gradients, axis=1, keepdims=True)
    else:
        raise ValueError('Unsupported loss function: {}'.format(loss_func))

def update(self, w, x, y, alpha, loss_func):
    """
    Updates the weight vector based on mean squared or half mean squared
    ↪loss

    Arg(s):
        w : numpy[float32]
            d x 1 weight vector
        x : numpy[float32]
            d x N feature vector
        y : numpy[float32]
            1 x N groundtruth vector
        alpha : float
            learning rate
        loss_func : str
            loss type either 'mean_squared', or 'half_mean_squared'
    Returns:
        numpy[float32] : d x 1 weights
    """

    # DONE: Implement the optimizer update function

    return w - alpha * self.__compute_gradients(w, x, y, loss_func)

```

Implementation of Linear Regression with Gradient Descent optimizer

```

[3]: class LinearRegressionGradientDescent(object):

    def __init__(self):
        # Define private variables
        self.__weights = None
        self.__optimizer = GradientDescentOptimizer()

    def fit(self, x, y, T, alpha, loss_func='mean_squared'):

```

```

'''
Fits the model to x and y by updating the weight vector
using gradient descent

Arg(s):
    x : numpy[float32]
        d x N feature vector
    y : numpy[float32]
        1 x N groundtruth vector
    T : int
        number of iterations to train
    alpha : float
        learning rate
    loss_func : str
        loss function to use
'''

# DONE: Implement the fit function
self.__weights = np.zeros([x.shape[0], 1])

for t in range(1, T + 1):

    # DONE: Compute loss function
    loss = self.__compute_loss(
        x=x,
        y=y,
        loss_func=loss_func)

    if (t % 10000) == 0:
        print('Step={} Loss={:.4f}'.format(t, loss))

    # DONE: Update weights
    self.__weights = self.__optimizer.update(
        w=self.__weights,
        x=x,
        y=y,
        alpha=alpha,
        loss_func=loss_func)

def predict(self, x):
    '''
    Predicts the label for each feature vector x

    Arg(s):
        x : numpy[float32]
            d x N feature vector
    Returns:

```

```

        numpy[float32] : 1 x N vector
    """

    # DONE: Implements the predict function

    return np.matmul(self.__weights.T, x)

def __compute_loss(self, x, y, loss_func):
    """
    Returns the gradient of the mean squared or half mean squared loss

    Arg(s):
        x : numpy[float32]
            d x N feature vector
        y : numpy[float32]
            1 x N groundtruth vector
        loss_func : str
            loss type either 'mean_squared', or 'half_mean_squared'
    Returns:
        float : loss
    """

    # DONE: Implements the __compute_loss function
    predictions = self.predict(x)

    if loss_func == 'mean_squared':
        # DONE: Implements loss for mean squared loss
        loss = np.mean((predictions - y) ** 2)
    elif loss_func == 'half_mean_squared':
        # DONE: Implements loss for half mean squared loss
        loss = 0.50 * np.mean((predictions - y) ** 2)
    else:
        raise ValueError('Unsupported loss function: {}'.format(loss_func))

    return loss

```

Implementing training and validation loop for linear regression

```

[4]: # Load Diabetes and California housing prices dataset
datasets = [
    skdata.load_diabetes(),
    skdata.fetch_california_housing()
]
dataset_names = [
    'Diabetes',
    'California housing prices'
]

```

```

# Loss functions to minimize
dataset_loss_funcs = [
    ['mean_squared', 'half_mean_squared'],
    ['mean_squared', 'half_mean_squared']
]

# TODO: Select learning rates (alpha) for mean squared and half mean squared
↳loss
dataset_alphas = [
    [1e-4, 2e-4],
    [1e-7, 2e-7]
]

# TODO: Select number of steps (T) to train for mean squared and half mean
↳squared loss
dataset_Ts = [
    [30000, 30000],
    [2500000, 2500000]
]

for dataset_options in zip(datasets, dataset_names, dataset_loss_funcs,
↳dataset_alphas, dataset_Ts):

    dataset, dataset_name, loss_funcs, alphas, Ts = dataset_options

    '''
    Create the training, validation and testing splits
    '''

    x = dataset.data
    y = dataset.target

    # Shuffle the dataset based on sample indices
    shuffled_indices = np.random.permutation(x.shape[0])

    # Choose the first 80% as training set, next 10% as validation and the rest
↳as testing
    train_split_idx = int(0.80 * x.shape[0])
    val_split_idx = int(0.90 * x.shape[0])

    train_indices = shuffled_indices[0:train_split_idx]
    val_indices = shuffled_indices[train_split_idx:val_split_idx]
    test_indices = shuffled_indices[val_split_idx:]

    # Select the examples from x and y to construct our training, validation,
↳testing sets
    x_train, y_train = x[train_indices, :], y[train_indices]

```

```

x_val, y_val = x[val_indices, :], y[val_indices]
x_test, y_test = x[test_indices, :], y[test_indices]

'''
Trains and tests Linear Regression model from scikit-learn
'''

# DONE: Initialize scikit-learn linear regression model without bias
model_scikit = LinearRegressionSciKit(fit_intercept=False)

# DONE: Trains scikit-learn linear regression model
model_scikit.fit(x_train, y_train)

print('***** Results of scikit-learn linear regression model on {} dataset_
↳*****'.format(
    dataset_name))

# DONE: Test model on training set
predictions_train = model_scikit.predict(x_train)

score_mse_train = skmetrics.mean_squared_error(y_train, predictions_train)
print('Training set mean squared error: {:.4f}'.format(score_mse_train))

score_r2_train = skmetrics.r2_score(y_train, predictions_train)
print('Training set r-squared scores: {:.4f}'.format(score_r2_train))

# DONE: Test model on validation set
predictions_val = model_scikit.predict(x_val)

score_mse_val = skmetrics.mean_squared_error(y_val, predictions_val)
print('Validation set mean squared error: {:.4f}'.format(score_mse_val))

score_r2_val = skmetrics.r2_score(y_val, predictions_val)
print('Validation set r-squared scores: {:.4f}'.format(score_r2_val))

# DONE: Test model on testing set
predictions_test = model_scikit.predict(x_test)

score_mse_test = skmetrics.mean_squared_error(y_test, predictions_test)
print('Testing set mean squared error: {:.4f}'.format(score_mse_test))

score_r2_test = skmetrics.r2_score(y_test, predictions_test)
print('Testing set r-squared scores: {:.4f}'.format(score_r2_test))

'''
Trains and tests our linear regression model using different solvers
'''

```



```

    # Take the transpose of the dataset to match the dimensions discussed in
    ↪lecture
    # i.e., (N x d) to (d x N)
    x_train = np.transpose(x_train, axes=(1, 0))
    x_val = np.transpose(x_val, axes=(1, 0))
    x_test = np.transpose(x_test, axes=(1, 0))
    y_train = np.expand_dims(y_train, axis=0)
    y_val = np.expand_dims(y_val, axis=0)
    y_test = np.expand_dims(y_test, axis=0)

    for loss_func, alpha, T in zip(loss_funcs, alphas, Ts):

        # DONE: Initialize our linear regression model
        model_ours = LinearRegressionGradientDescent()

        print('***** Results of our linear regression model trained with {}
    ↪loss, alpha={} and T={} on {} dataset *****'.format(
            loss_func, alpha, T, dataset_name))

        # DONE: Train model on training set
        model_ours.fit(
            x=x_train,
            y=y_train,
            T=T,
            alpha=alpha,
            loss_func=loss_func)

        # DONE: Test model on training set using mean squared error and
    ↪r-squared
        predictions_train = model_ours.predict(x_train)

        # Squeeze to remove extra dimensions before passing to r2_score
        y_train_squeezed = np.squeeze(y_train)
        predictions_train_squeezed = np.squeeze(predictions_train)

        score_mse_train = skmetrics.mean_squared_error(y_train_squeezed,
    ↪predictions_train_squeezed)
        print('Training set mean squared error: {:.4f}'.format(score_mse_train))

        score_r2_train = skmetrics.r2_score(y_train_squeezed,
    ↪predictions_train_squeezed)
        print('Training set r-squared scores: {:.4f}'.format(score_r2_train))

        # DONE: Test model on validation set using mean squared error and
    ↪r-squared

```

```

predictions_val = model_ours.predict(x_val)

# Squeeze to remove extra dimensions
y_val_squeezed = np.squeeze(y_val)
predictions_val_squeezed = np.squeeze(predictions_val)

score_mse_val = skmetrics.mean_squared_error(y_val_squeezed,
↪predictions_val_squeezed)
print('Validation set mean squared error: {:.4f}'.format(score_mse_val))

score_r2_val = skmetrics.r2_score(y_val_squeezed,
↪predictions_val_squeezed)
print('Validation set r-squared scores: {:.4f}'.format(score_r2_val))

# DONE: Test model on testing set using mean squared error and r-squared
predictions_test = model_ours.predict(x_test)

# Squeeze to remove extra dimensions
y_test_squeezed = np.squeeze(y_test)
predictions_test_squeezed = np.squeeze(predictions_test)

score_mse_test = skmetrics.mean_squared_error(y_test_squeezed,
↪predictions_test_squeezed)
print('Testing set mean squared error: {:.4f}'.format(score_mse_test))

score_r2_test = skmetrics.r2_score(y_test_squeezed,
↪predictions_test_squeezed)
print('Testing set r-squared scores: {:.4f}'.format(score_r2_test))

```

***** Results of scikit-learn linear regression model on Diabetes dataset *****

Training set mean squared error: 25472.6663

Training set r-squared scores: -3.2907

Validation set mean squared error: 30376.7403

Validation set r-squared scores: -5.1176

Testing set mean squared error: 27024.0792

Testing set r-squared scores: -3.2604

***** Results of our linear regression model trained with mean_squared loss,
alpha=0.0001 and T=30000 on Diabetes dataset *****

Step=10000 Loss=28180.3086

Step=20000 Loss=28106.3236

Step=30000 Loss=28034.8016

Training set mean squared error: 28034.7946

Training set r-squared scores: -3.7223

Validation set mean squared error: 34538.4871

Validation set r-squared scores: -5.9557

Testing set mean squared error: 29713.4975

Testing set r-squared scores: -3.6844

***** Results of our linear regression model trained with half_mean_squared loss, alpha=0.0002 and T=30000 on Diabetes dataset *****

Step=10000 Loss=14090.1543

Step=20000 Loss=14053.1618

Step=30000 Loss=14017.4008

Training set mean squared error: 28034.7946

Training set r-squared scores: -3.7223

Validation set mean squared error: 34538.4871

Validation set r-squared scores: -5.9557

Testing set mean squared error: 29713.4975

Testing set r-squared scores: -3.6844

***** Results of scikit-learn linear regression model on California housing prices dataset *****

Training set mean squared error: 0.5989

Training set r-squared scores: 0.5463

Validation set mean squared error: 0.5922

Validation set r-squared scores: 0.5721

Testing set mean squared error: 0.6590

Testing set r-squared scores: 0.5185

***** Results of our linear regression model trained with mean_squared loss, alpha=1e-07 and T=2500000 on California housing prices dataset *****

Step=10000 Loss=1.2972

Step=20000 Loss=1.2855

Step=30000 Loss=1.2744

Step=40000 Loss=1.2638

Step=50000 Loss=1.2536

Step=60000 Loss=1.2437

Step=70000 Loss=1.2340

Step=80000 Loss=1.2245

Step=90000 Loss=1.2151

Step=100000 Loss=1.2060

Step=110000 Loss=1.1970

Step=120000 Loss=1.1882

Step=130000 Loss=1.1796

Step=140000 Loss=1.1711

Step=150000 Loss=1.1628

Step=160000 Loss=1.1546

Step=170000 Loss=1.1465

Step=180000 Loss=1.1386

Step=190000 Loss=1.1308

Step=200000 Loss=1.1232

Step=210000 Loss=1.1157

Step=220000 Loss=1.1083

Step=230000 Loss=1.1010

Step=240000 Loss=1.0939

Step=250000 Loss=1.0869

Step=260000 Loss=1.0800

Step=270000 Loss=1.0732

Step=280000	Loss=1.0666
Step=290000	Loss=1.0600
Step=300000	Loss=1.0536
Step=310000	Loss=1.0473
Step=320000	Loss=1.0410
Step=330000	Loss=1.0349
Step=340000	Loss=1.0289
Step=350000	Loss=1.0230
Step=360000	Loss=1.0172
Step=370000	Loss=1.0114
Step=380000	Loss=1.0058
Step=390000	Loss=1.0003
Step=400000	Loss=0.9948
Step=410000	Loss=0.9895
Step=420000	Loss=0.9842
Step=430000	Loss=0.9790
Step=440000	Loss=0.9739
Step=450000	Loss=0.9689
Step=460000	Loss=0.9640
Step=470000	Loss=0.9591
Step=480000	Loss=0.9543
Step=490000	Loss=0.9496
Step=500000	Loss=0.9450
Step=510000	Loss=0.9404
Step=520000	Loss=0.9360
Step=530000	Loss=0.9315
Step=540000	Loss=0.9272
Step=550000	Loss=0.9229
Step=560000	Loss=0.9187
Step=570000	Loss=0.9146
Step=580000	Loss=0.9105
Step=590000	Loss=0.9065
Step=600000	Loss=0.9026
Step=610000	Loss=0.8987
Step=620000	Loss=0.8948
Step=630000	Loss=0.8911
Step=640000	Loss=0.8874
Step=650000	Loss=0.8837
Step=660000	Loss=0.8801
Step=670000	Loss=0.8766
Step=680000	Loss=0.8731
Step=690000	Loss=0.8697
Step=700000	Loss=0.8663
Step=710000	Loss=0.8630
Step=720000	Loss=0.8597
Step=730000	Loss=0.8564
Step=740000	Loss=0.8533
Step=750000	Loss=0.8501

Step=760000	Loss=0.8471
Step=770000	Loss=0.8440
Step=780000	Loss=0.8410
Step=790000	Loss=0.8381
Step=800000	Loss=0.8352
Step=810000	Loss=0.8323
Step=820000	Loss=0.8295
Step=830000	Loss=0.8267
Step=840000	Loss=0.8240
Step=850000	Loss=0.8213
Step=860000	Loss=0.8186
Step=870000	Loss=0.8160
Step=880000	Loss=0.8135
Step=890000	Loss=0.8109
Step=900000	Loss=0.8084
Step=910000	Loss=0.8059
Step=920000	Loss=0.8035
Step=930000	Loss=0.8011
Step=940000	Loss=0.7988
Step=950000	Loss=0.7964
Step=960000	Loss=0.7942
Step=970000	Loss=0.7919
Step=980000	Loss=0.7897
Step=990000	Loss=0.7875
Step=1000000	Loss=0.7853
Step=1010000	Loss=0.7832
Step=1020000	Loss=0.7811
Step=1030000	Loss=0.7790
Step=1040000	Loss=0.7770
Step=1050000	Loss=0.7750
Step=1060000	Loss=0.7730
Step=1070000	Loss=0.7710
Step=1080000	Loss=0.7691
Step=1090000	Loss=0.7672
Step=1100000	Loss=0.7654
Step=1110000	Loss=0.7635
Step=1120000	Loss=0.7617
Step=1130000	Loss=0.7599
Step=1140000	Loss=0.7581
Step=1150000	Loss=0.7564
Step=1160000	Loss=0.7547
Step=1170000	Loss=0.7530
Step=1180000	Loss=0.7513
Step=1190000	Loss=0.7497
Step=1200000	Loss=0.7480
Step=1210000	Loss=0.7464
Step=1220000	Loss=0.7449
Step=1230000	Loss=0.7433

Step=1240000	Loss=0.7418
Step=1250000	Loss=0.7403
Step=1260000	Loss=0.7388
Step=1270000	Loss=0.7373
Step=1280000	Loss=0.7359
Step=1290000	Loss=0.7344
Step=1300000	Loss=0.7330
Step=1310000	Loss=0.7316
Step=1320000	Loss=0.7302
Step=1330000	Loss=0.7289
Step=1340000	Loss=0.7276
Step=1350000	Loss=0.7262
Step=1360000	Loss=0.7249
Step=1370000	Loss=0.7237
Step=1380000	Loss=0.7224
Step=1390000	Loss=0.7212
Step=1400000	Loss=0.7199
Step=1410000	Loss=0.7187
Step=1420000	Loss=0.7175
Step=1430000	Loss=0.7163
Step=1440000	Loss=0.7152
Step=1450000	Loss=0.7140
Step=1460000	Loss=0.7129
Step=1470000	Loss=0.7118
Step=1480000	Loss=0.7107
Step=1490000	Loss=0.7096
Step=1500000	Loss=0.7085
Step=1510000	Loss=0.7075
Step=1520000	Loss=0.7064
Step=1530000	Loss=0.7054
Step=1540000	Loss=0.7044
Step=1550000	Loss=0.7034
Step=1560000	Loss=0.7024
Step=1570000	Loss=0.7014
Step=1580000	Loss=0.7005
Step=1590000	Loss=0.6995
Step=1600000	Loss=0.6986
Step=1610000	Loss=0.6976
Step=1620000	Loss=0.6967
Step=1630000	Loss=0.6958
Step=1640000	Loss=0.6949
Step=1650000	Loss=0.6941
Step=1660000	Loss=0.6932
Step=1670000	Loss=0.6924
Step=1680000	Loss=0.6915
Step=1690000	Loss=0.6907
Step=1700000	Loss=0.6899
Step=1710000	Loss=0.6891

Step=1720000	Loss=0.6883
Step=1730000	Loss=0.6875
Step=1740000	Loss=0.6867
Step=1750000	Loss=0.6859
Step=1760000	Loss=0.6852
Step=1770000	Loss=0.6844
Step=1780000	Loss=0.6837
Step=1790000	Loss=0.6830
Step=1800000	Loss=0.6823
Step=1810000	Loss=0.6815
Step=1820000	Loss=0.6809
Step=1830000	Loss=0.6802
Step=1840000	Loss=0.6795
Step=1850000	Loss=0.6788
Step=1860000	Loss=0.6781
Step=1870000	Loss=0.6775
Step=1880000	Loss=0.6769
Step=1890000	Loss=0.6762
Step=1900000	Loss=0.6756
Step=1910000	Loss=0.6750
Step=1920000	Loss=0.6744
Step=1930000	Loss=0.6738
Step=1940000	Loss=0.6732
Step=1950000	Loss=0.6726
Step=1960000	Loss=0.6720
Step=1970000	Loss=0.6714
Step=1980000	Loss=0.6708
Step=1990000	Loss=0.6703
Step=2000000	Loss=0.6697
Step=2010000	Loss=0.6692
Step=2020000	Loss=0.6687
Step=2030000	Loss=0.6681
Step=2040000	Loss=0.6676
Step=2050000	Loss=0.6671
Step=2060000	Loss=0.6666
Step=2070000	Loss=0.6661
Step=2080000	Loss=0.6656
Step=2090000	Loss=0.6651
Step=2100000	Loss=0.6646
Step=2110000	Loss=0.6641
Step=2120000	Loss=0.6637
Step=2130000	Loss=0.6632
Step=2140000	Loss=0.6627
Step=2150000	Loss=0.6623
Step=2160000	Loss=0.6618
Step=2170000	Loss=0.6614
Step=2180000	Loss=0.6610
Step=2190000	Loss=0.6605

```

Step=2200000 Loss=0.6601
Step=2210000 Loss=0.6597
Step=2220000 Loss=0.6593
Step=2230000 Loss=0.6589
Step=2240000 Loss=0.6585
Step=2250000 Loss=0.6581
Step=2260000 Loss=0.6577
Step=2270000 Loss=0.6573
Step=2280000 Loss=0.6569
Step=2290000 Loss=0.6565
Step=2300000 Loss=0.6561
Step=2310000 Loss=0.6558
Step=2320000 Loss=0.6554
Step=2330000 Loss=0.6550
Step=2340000 Loss=0.6547
Step=2350000 Loss=0.6543
Step=2360000 Loss=0.6540
Step=2370000 Loss=0.6536
Step=2380000 Loss=0.6533
Step=2390000 Loss=0.6530
Step=2400000 Loss=0.6526
Step=2410000 Loss=0.6523
Step=2420000 Loss=0.6520
Step=2430000 Loss=0.6517
Step=2440000 Loss=0.6514
Step=2450000 Loss=0.6511
Step=2460000 Loss=0.6508
Step=2470000 Loss=0.6505
Step=2480000 Loss=0.6502
Step=2490000 Loss=0.6499
Step=2500000 Loss=0.6496
Training set mean squared error: 0.6496
Training set r-squared scores: 0.5080
Validation set mean squared error: 0.6348
Validation set r-squared scores: 0.5413
Testing set mean squared error: 0.7028
Testing set r-squared scores: 0.4865
***** Results of our linear regression model trained with half_mean_squared
loss, alpha=2e-07 and T=2500000 on California housing prices dataset *****
Step=10000 Loss=0.6486
Step=20000 Loss=0.6427
Step=30000 Loss=0.6372
Step=40000 Loss=0.6319
Step=50000 Loss=0.6268
Step=60000 Loss=0.6218
Step=70000 Loss=0.6170
Step=80000 Loss=0.6122
Step=90000 Loss=0.6076

```


Step=100000	Loss=0.6030
Step=110000	Loss=0.5985
Step=120000	Loss=0.5941
Step=130000	Loss=0.5898
Step=140000	Loss=0.5855
Step=150000	Loss=0.5814
Step=160000	Loss=0.5773
Step=170000	Loss=0.5733
Step=180000	Loss=0.5693
Step=190000	Loss=0.5654
Step=200000	Loss=0.5616
Step=210000	Loss=0.5578
Step=220000	Loss=0.5541
Step=230000	Loss=0.5505
Step=240000	Loss=0.5469
Step=250000	Loss=0.5434
Step=260000	Loss=0.5400
Step=270000	Loss=0.5366
Step=280000	Loss=0.5333
Step=290000	Loss=0.5300
Step=300000	Loss=0.5268
Step=310000	Loss=0.5236
Step=320000	Loss=0.5205
Step=330000	Loss=0.5175
Step=340000	Loss=0.5144
Step=350000	Loss=0.5115
Step=360000	Loss=0.5086
Step=370000	Loss=0.5057
Step=380000	Loss=0.5029
Step=390000	Loss=0.5001
Step=400000	Loss=0.4974
Step=410000	Loss=0.4947
Step=420000	Loss=0.4921
Step=430000	Loss=0.4895
Step=440000	Loss=0.4870
Step=450000	Loss=0.4844
Step=460000	Loss=0.4820
Step=470000	Loss=0.4795
Step=480000	Loss=0.4772
Step=490000	Loss=0.4748
Step=500000	Loss=0.4725
Step=510000	Loss=0.4702
Step=520000	Loss=0.4680
Step=530000	Loss=0.4658
Step=540000	Loss=0.4636
Step=550000	Loss=0.4615
Step=560000	Loss=0.4594
Step=570000	Loss=0.4573

Step=580000	Loss=0.4553
Step=590000	Loss=0.4532
Step=600000	Loss=0.4513
Step=610000	Loss=0.4493
Step=620000	Loss=0.4474
Step=630000	Loss=0.4455
Step=640000	Loss=0.4437
Step=650000	Loss=0.4419
Step=660000	Loss=0.4401
Step=670000	Loss=0.4383
Step=680000	Loss=0.4365
Step=690000	Loss=0.4348
Step=700000	Loss=0.4331
Step=710000	Loss=0.4315
Step=720000	Loss=0.4298
Step=730000	Loss=0.4282
Step=740000	Loss=0.4266
Step=750000	Loss=0.4251
Step=760000	Loss=0.4235
Step=770000	Loss=0.4220
Step=780000	Loss=0.4205
Step=790000	Loss=0.4190
Step=800000	Loss=0.4176
Step=810000	Loss=0.4162
Step=820000	Loss=0.4147
Step=830000	Loss=0.4134
Step=840000	Loss=0.4120
Step=850000	Loss=0.4106
Step=860000	Loss=0.4093
Step=870000	Loss=0.4080
Step=880000	Loss=0.4067
Step=890000	Loss=0.4055
Step=900000	Loss=0.4042
Step=910000	Loss=0.4030
Step=920000	Loss=0.4018
Step=930000	Loss=0.4006
Step=940000	Loss=0.3994
Step=950000	Loss=0.3982
Step=960000	Loss=0.3971
Step=970000	Loss=0.3959
Step=980000	Loss=0.3948
Step=990000	Loss=0.3937
Step=1000000	Loss=0.3927
Step=1010000	Loss=0.3916
Step=1020000	Loss=0.3905
Step=1030000	Loss=0.3895
Step=1040000	Loss=0.3885
Step=1050000	Loss=0.3875

Step=1060000	Loss=0.3865
Step=1070000	Loss=0.3855
Step=1080000	Loss=0.3846
Step=1090000	Loss=0.3836
Step=1100000	Loss=0.3827
Step=1110000	Loss=0.3818
Step=1120000	Loss=0.3808
Step=1130000	Loss=0.3800
Step=1140000	Loss=0.3791
Step=1150000	Loss=0.3782
Step=1160000	Loss=0.3773
Step=1170000	Loss=0.3765
Step=1180000	Loss=0.3757
Step=1190000	Loss=0.3748
Step=1200000	Loss=0.3740
Step=1210000	Loss=0.3732
Step=1220000	Loss=0.3724
Step=1230000	Loss=0.3717
Step=1240000	Loss=0.3709
Step=1250000	Loss=0.3701
Step=1260000	Loss=0.3694
Step=1270000	Loss=0.3687
Step=1280000	Loss=0.3679
Step=1290000	Loss=0.3672
Step=1300000	Loss=0.3665
Step=1310000	Loss=0.3658
Step=1320000	Loss=0.3651
Step=1330000	Loss=0.3644
Step=1340000	Loss=0.3638
Step=1350000	Loss=0.3631
Step=1360000	Loss=0.3625
Step=1370000	Loss=0.3618
Step=1380000	Loss=0.3612
Step=1390000	Loss=0.3606
Step=1400000	Loss=0.3600
Step=1410000	Loss=0.3594
Step=1420000	Loss=0.3588
Step=1430000	Loss=0.3582
Step=1440000	Loss=0.3576
Step=1450000	Loss=0.3570
Step=1460000	Loss=0.3564
Step=1470000	Loss=0.3559
Step=1480000	Loss=0.3553
Step=1490000	Loss=0.3548
Step=1500000	Loss=0.3543
Step=1510000	Loss=0.3537
Step=1520000	Loss=0.3532
Step=1530000	Loss=0.3527

Step=1540000	Loss=0.3522
Step=1550000	Loss=0.3517
Step=1560000	Loss=0.3512
Step=1570000	Loss=0.3507
Step=1580000	Loss=0.3502
Step=1590000	Loss=0.3498
Step=1600000	Loss=0.3493
Step=1610000	Loss=0.3488
Step=1620000	Loss=0.3484
Step=1630000	Loss=0.3479
Step=1640000	Loss=0.3475
Step=1650000	Loss=0.3470
Step=1660000	Loss=0.3466
Step=1670000	Loss=0.3462
Step=1680000	Loss=0.3458
Step=1690000	Loss=0.3453
Step=1700000	Loss=0.3449
Step=1710000	Loss=0.3445
Step=1720000	Loss=0.3441
Step=1730000	Loss=0.3437
Step=1740000	Loss=0.3434
Step=1750000	Loss=0.3430
Step=1760000	Loss=0.3426
Step=1770000	Loss=0.3422
Step=1780000	Loss=0.3418
Step=1790000	Loss=0.3415
Step=1800000	Loss=0.3411
Step=1810000	Loss=0.3408
Step=1820000	Loss=0.3404
Step=1830000	Loss=0.3401
Step=1840000	Loss=0.3397
Step=1850000	Loss=0.3394
Step=1860000	Loss=0.3391
Step=1870000	Loss=0.3387
Step=1880000	Loss=0.3384
Step=1890000	Loss=0.3381
Step=1900000	Loss=0.3378
Step=1910000	Loss=0.3375
Step=1920000	Loss=0.3372
Step=1930000	Loss=0.3369
Step=1940000	Loss=0.3366
Step=1950000	Loss=0.3363
Step=1960000	Loss=0.3360
Step=1970000	Loss=0.3357
Step=1980000	Loss=0.3354
Step=1990000	Loss=0.3351
Step=2000000	Loss=0.3349
Step=2010000	Loss=0.3346

Step=2020000	Loss=0.3343
Step=2030000	Loss=0.3341
Step=2040000	Loss=0.3338
Step=2050000	Loss=0.3335
Step=2060000	Loss=0.3333
Step=2070000	Loss=0.3330
Step=2080000	Loss=0.3328
Step=2090000	Loss=0.3325
Step=2100000	Loss=0.3323
Step=2110000	Loss=0.3321
Step=2120000	Loss=0.3318
Step=2130000	Loss=0.3316
Step=2140000	Loss=0.3314
Step=2150000	Loss=0.3311
Step=2160000	Loss=0.3309
Step=2170000	Loss=0.3307
Step=2180000	Loss=0.3305
Step=2190000	Loss=0.3303
Step=2200000	Loss=0.3300
Step=2210000	Loss=0.3298
Step=2220000	Loss=0.3296
Step=2230000	Loss=0.3294
Step=2240000	Loss=0.3292
Step=2250000	Loss=0.3290
Step=2260000	Loss=0.3288
Step=2270000	Loss=0.3286
Step=2280000	Loss=0.3284
Step=2290000	Loss=0.3283
Step=2300000	Loss=0.3281
Step=2310000	Loss=0.3279
Step=2320000	Loss=0.3277
Step=2330000	Loss=0.3275
Step=2340000	Loss=0.3273
Step=2350000	Loss=0.3272
Step=2360000	Loss=0.3270
Step=2370000	Loss=0.3268
Step=2380000	Loss=0.3267
Step=2390000	Loss=0.3265
Step=2400000	Loss=0.3263
Step=2410000	Loss=0.3262
Step=2420000	Loss=0.3260
Step=2430000	Loss=0.3258
Step=2440000	Loss=0.3257
Step=2450000	Loss=0.3255
Step=2460000	Loss=0.3254
Step=2470000	Loss=0.3252
Step=2480000	Loss=0.3251
Step=2490000	Loss=0.3249

```
Step=2500000  Loss=0.3248
Training set mean squared error: 0.6496
Training set r-squared scores: 0.5080
Validation set mean squared error: 0.6348
Validation set r-squared scores: 0.5413
Testing set mean squared error: 0.7028
Testing set r-squared scores: 0.4865
```

[]: